All-vertex survey: first 12 audited strata

Run 20261017-survey-prefix-12. Frozen checkpoint, not completion.

First 12 strata: projected size-64 MITM relation yield

Two public targets per vertex; mean with observed min–max. All strata are separate graph/stratum pairs.

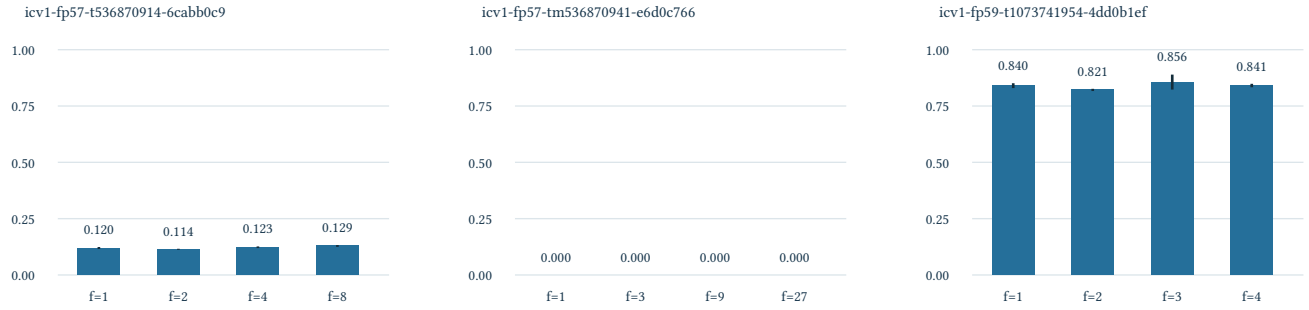

f = conductor of End(E). One vertex per displayed stratum; variation is not a graph-wide distribution.

Native recorded coverage: 12 audited vertices in 12 distinct graph/stratum pairs. Four vertices in each graph. The immutable queue contains all 24,823 vertices and 51 strata; the first phase visits one vertex per stratum before all remaining vertices. Pilot/root reruns use the same targets and seeds; they are repeated observations, not additional independent targets.

Observed outcomes: 192 measured attempts: 119 verified, 58 exhausted, 15 watchdog timeouts. All 12 standard audits pass. Each successful deterministic record replays, and each exhausted deterministic record passes the additional native replay receipt. No wrong answers were recorded. Timeouts have unknown final relation and operation counters. Their elapsed observations are retained.

Recorded direct-IC counters: 301,715 attempted decompositions, 6,368 relations found across the 120 direct-IC attempts with final reports. Numerator and denominator include exhausted reports and the raw-base control. They are not 120 successful DLPs or independent identically distributed trials. Phase, rank, dependency, failure, memory and unpriced-work records remain available.

Subset below: size-64 projected MITM only. Cold GAE are lower bounds with unpriced native work; no fully priced cost or runtime ranking follows. Each row is two measured one-target attempts on one exact vertex.

Graph key	f	Status	n	GAE min	GAE mean	GAE max	Mean yield
two	1	verified	2	19,530	19,970.0	20,410	0.119769
two	2	verified	2	19,505	19,546.5	19,588	0.114059
two	4	verified	2	19,466	19,531.0	19,596	0.123116
two	8	verified	2	19,467	19,537.5	19,608	0.128813
three	1	exhausted	2	20,239	20,252.0	20,265	0.000000
three	3	exhausted	2	20,245	20,268.0	20,291	0.000000
three	9	exhausted	2	20,249	20,251.0	20,253	0.000000
three	27	exhausted	2	20,249	20,259.5	20,270	0.000000
mixed	1	verified	2	20,463	20,467.5	20,472	0.840387
mixed	2	verified	2	20,216	20,228.5	20,241	0.821442
mixed	3	verified	2	20,281	20,438.0	20,595	0.856216
mixed	4	verified	2	20,095	20,123.0	20,151	0.840969

Conditions: actual largest-prime subgroups; projected sizes 16/64 with MITM/subtraction, unchanged raw-base MITM, rho.negation, textbook BSGS and rho A/A. Two public targets, one round, no warmup, 4,096 trials, 3-second watchdog. Prime-subgroup membership and all recovered scalars are checked. Every run charges its cold solver setup; no reusable logs are assumed free. Ordinary finite-field IC remains unsupported at the actual recorded embedding degrees.

Host: Intel Xeon Platinum 8573C, Linux x86-64, four-CPU quota, 16 GiB RAM. Unoptimized ecbench SHA-256 1869ca2e7be271794df8128c08a6dca0201038e2c81a3e0ce6434f7ec0d8379d; native driver source commit ee54ad603. All records earn L0; raw solve/process times are descriptive, not admitted performance. Short driver rejection checks ran during this counts-only block; no build or test suite ran during it. The observed invocation took 272 elapsed seconds, including runner overhead and audits. This is a local scheduling observation, not an all-vertex forecast.

Evidence: execution/campaigns/survey-prefix-12, immutable full queue in the delivery archive, campaign.json hashes, SQLite checkpoint, standard audits and exhaustion-replay receipts. The native summary v2 includes queued/audited coverage by conductor and explicit sample counts for every min/max/mean. Common 0/0.1/1 GAE pricing-band output is separate. Same-implementation replay is not independent implementation validation. A second writer and changed driver were rejected; zero-work resume preserved the 12 audited sessions.

The original full matrix (million trials, eight targets, five rounds, warmup, four base sizes), original A–E full graphs, exact representation registration, transported targets, full composite PH, qualified runtimes, finite-field IC and independent implementation replay remain open. A real queue is not a completed solve. Canonical best-method/frontier claims are unchanged.