Boolean matrix step: F4 and F5

One shared K_1 / 2^23 m=2 system, 22 Boolean variables, maximum degree 4. Matrix construction, reduction and readback; F5 includes its criterion.

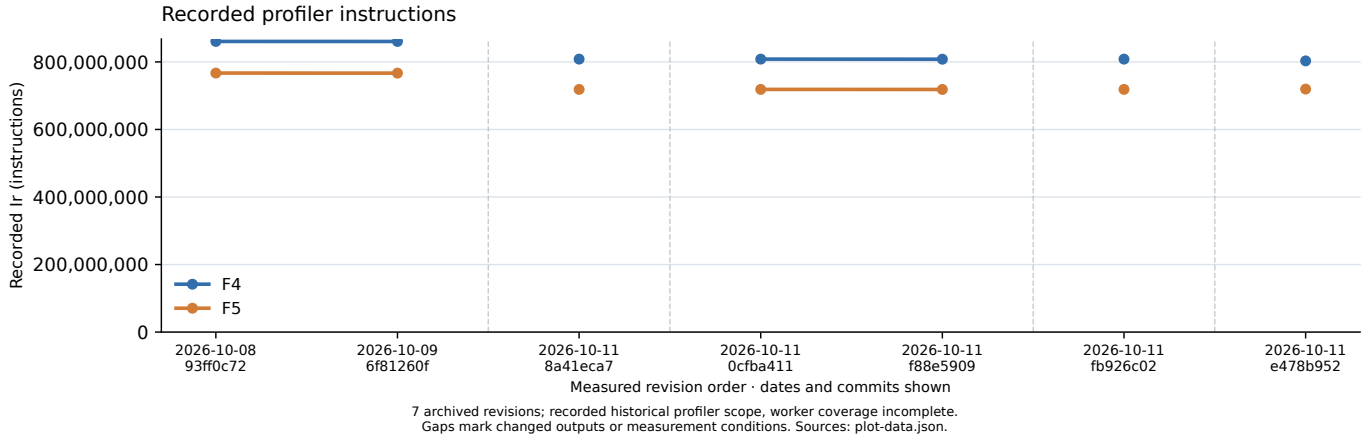

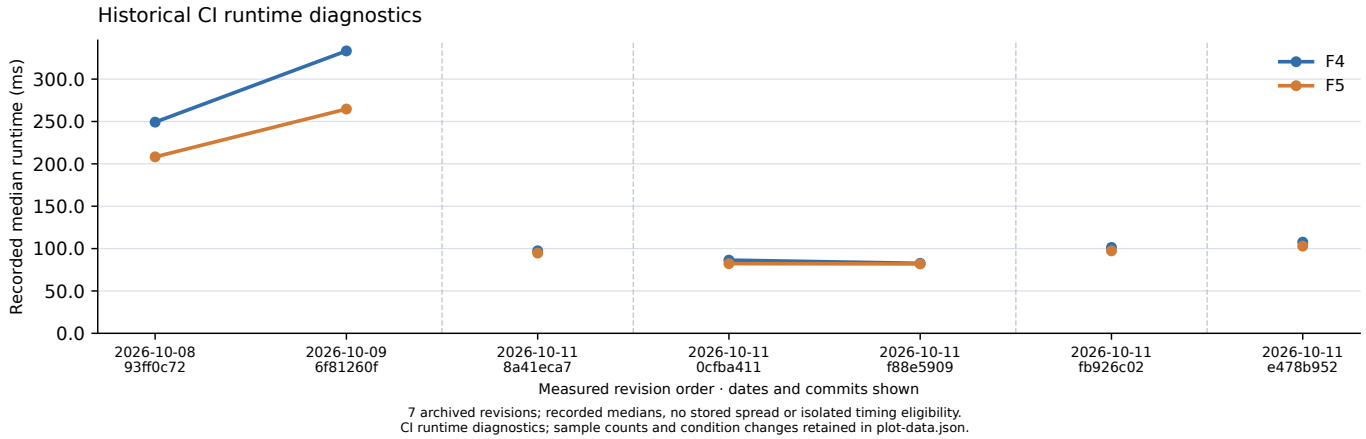

UTC date / revision	F4 Ir	F5 Ir	F4 ms	F5 ms
2026-10-08 / 93ff0c72	860,633,889	766,940,469	249.241	208.250
2026-10-09 / 6f81260f	860,555,567	766,776,087	333.279	264.687
2026-10-11 / 8a41eca7	808,263,705	718,576,235	97.380	94.658
2026-10-11 / 0cfba411	808,263,705	718,578,523	86.390	82.214
2026-10-11 / f88e5909	807,989,935	718,391,055	82.597	81.948
2026-10-11 / fb926c02	808,263,705	718,576,189	101.379	97.323
2026-10-11 / e478b952	803,017,657	719,516,063	107.585	102.927

Historical Python-orchestrated perf.history/v1 snapshots; no new performance run or speedup claim.

Runtime: one recorded median per kernel and revision; sample counts retained per point, no stored spread, no isolated timing eligibility. Runtime is diagnostic.

Profiler Ir is the recorded historical scope. Worker-thread coverage is incomplete for these kernels; it is not a complete solver operation count.

Lines stop at missing values, changed output fingerprints, or changed host/compiler/harness/profiler. Different algorithm fingerprints are expected and are not equated.

X positions are measured revision order, labelled with UTC commit date and revision; spacing is not elapsed time.

Evidence: 7 imported perf.history/v1 snapshots; historical_sources in plot-data.json retains raw filenames, SHA-256 hashes, hosts, harnesses and sample counts. Source data and provenance: plot-data.json and REPORT.md. These plots do not establish an end-to-end ECDLP gain.